

Deploying Applications in Kubernetes

Anthony E. Nocentino
aen@centinosystems.com



Course Overview

- **Module 0** - Introduction
- **Module 1** - Container Fundamentals and Introducing Kubernetes
- **Module 2** - Kubernetes Architecture and API Objects
- **Lunch @ 12:00-12:45**
- **Module 3** - Interacting With Your Cluster
- **Module 4** - Deploying Applications in Kubernetes
- **Module 5** - Building and Deploying Container-based Applications in Kubernetes

Agenda

- **Deploying Applications in Kubernetes**
 - Namespaces
 - Labels
 - Services
 - Deployments

Organizing Objects in Kubernetes

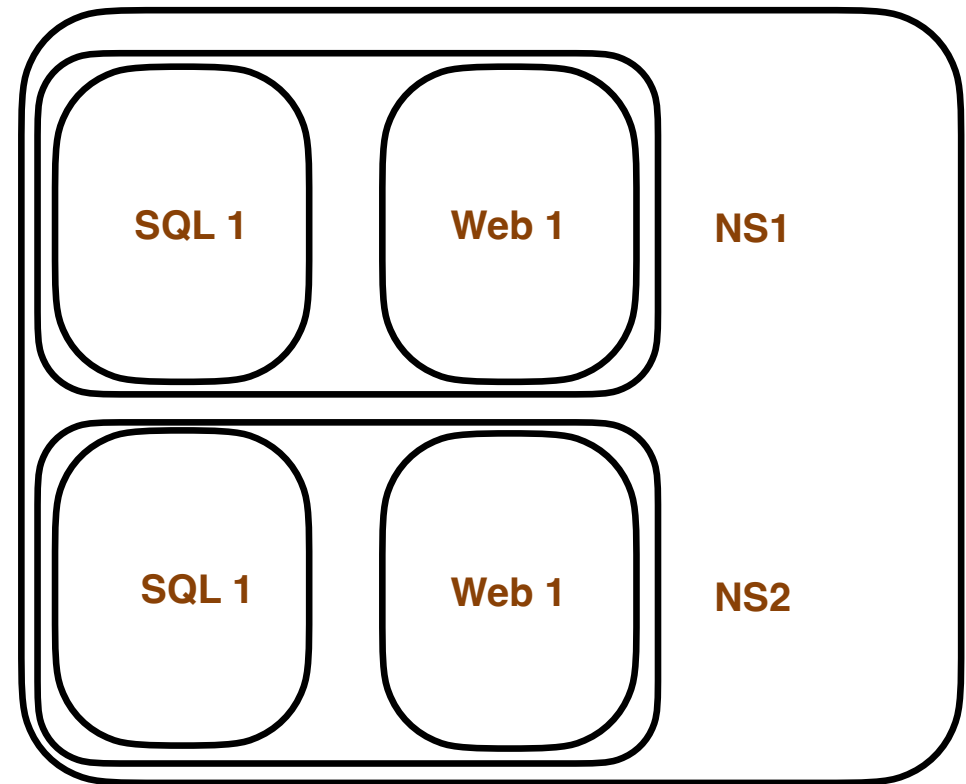
- **Namespaces**
- **Labels**
- **Annotations**

Namespaces

- Ability to subdivide a cluster and its resources
- Conceptually a “virtual cluster”
 - Resource isolation/organization
 - Security boundary for Role Based Access Controls
 - Naming boundary
- A resource can be in only one namespace
- Has nothing to do with the concept of a Linux namespace

Working with Namespaces


- Create/Query/Delete
- Operate on objects in a Namespace
- Some objects are Namespaced...
some aren't
- Resources are Namespaces...
Pods, Controllers, Services
- Physical things are not...
PersistentVolumes, Nodes



Cluster

Working with Namespaces

- **kube-public**
- **kube-system**
- User Defined
 - Imperatively with **kubectl**
 - Declaratively in a Manifest in YAML



Cluster

Labels

- Used to organize resources - Pods, Nodes and more
- Enables you to perform operations on a collection of resources...like Pods
- Influence internal operations of Kubernetes
- Non-hierarchical, key/value pair
- Have more than one label per resource
- Label Selectors
 - Used to select/query objects
 - Return collections of objects that satisfy search conditions

Working With Labels

- Creating resources with Labels
 - Imperatively with **kubect1**
 - Declaratively in a Manifest in YAML
- Editing existing resources' Labels
 - Assign a new Label
 - Overwriting an existing Label

Adding and Editing Labels

```
kubectl label pod nginx tier=PROD app=v1
```

```
kubectl label pod nginx tier=DEBUG app=v1 --overwrite
```

```
kubectl label pod nginx app-
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx-pod
```

```
  labels:
```

```
    tier: PROD
```

```
    app: v1
```

```
spec:
```

```
  containers:
```

```
  - name: nginx
```

```
...
```



Querying Using Labels and Selectors

```
kubectl get pods --show-labels
```

```
kubectl get pods --selector tier=prod
```

```
kubectl get pods -l tier=prod
```

```
kubectl get pods -l 'tier in (prod,qa)'
```

```
kubectl get pods -l 'tier notin (prod,qa)'
```

```
kubectl get nodes --show-labels
```


How Kubernetes Uses Labels

- **Deployments**, **ReplicaSets** and **Services** match Pods using selector labels
 - Pod Selector - **matchLabels** and **matchExpressions**
- Influencing Scheduling Pods based on labels
- Nodes (specifically for access to hardware, perhaps high speed disk or GPU, deploying around fault domains)
 - Specified Node name with **nodeAffinity**
 - Using a label selector with **nodeSelector**

Using Services

- Due to the ephemerality of Pods, ReplicaSets ability to add/remove Pods
- Provide persistent IP and DNS to Pod based applications
- Services use Label Selectors to “select” the Pods in a Service,
 - Registering the Endpoints in the Service (Pod IP and Port pair)
- Implemented in the kube-proxy on the Node in iptables (kernel mode)
- kube-proxy watches the API Server and the Endpoints

Using Services



Types of Services

- **NodePort** - available on the internal Cluster's IP/Port and on each Node
- **ClusterIP** - (default) - IP accessible only inside the cluster
- **LoadBalancer** - Cloud specific load balancer, redirects to Node services

Defining Services


```
kubectl expose deployment hello-world --port=80 --target-port=8080
```

```
apiVersion: v1
kind: Service
metadata:
  name: hello-world
spec:
  selector:
    app: hello-world
  ports:
    - port: 80
      protocol: TCP
      targetPort: 8080
```


Using Deployments

- **Deployments** are used to provide declarative updates to Pods and **ReplicaSets**
- We define the state and use the Deployment Controller to move towards that state
- **Deployments** are made of **ReplicaSets** and manage the transition between the **ReplicaSets**
- The pod-template hash Label is used to differentiate between versions of **ReplicaSets**

Controller Operations - Deployment



Using Deployments

- Give you the controlled rollout of a new version of your application
 - Control the rate of deploying new Pods
 - **Recreate** vs. **RollingUpdate**
 - Control rollout with health checks
 - **maxSurge**
 - **maxUnavailable**
- Rollback to an earlier version
- Scale based on load
- Pause to make corrections

Defining a Deployment



Hands on Lab

- Working with Namespaces and Labels
- Using Services and Endpoints
- Using Deployments
 - Rollout
 - Changing an image
 - Rolling back
- Understanding how Kubernetes uses Labels in Controllers and Services
 - **Deployments**
 - **ReplicaSets**

Review

- **Deploying Applications in Kubernetes**
 - Namespaces
 - Labels
 - Services
 - Deployments